

PAR – 2nd In-Term Exam – Course 2018/19-Q1
December 19th, 2018

Problem 1 (3 points) Given the following OpenMP parallel function that builds some sort of histogram `h` for the elements of vector `v` received as argument:

```
#define SIZE1 131072
#define SIZE2 128
#define MAX 1024
#define CACHE_LINE_SIZE 64 // size of the cache line in bytes
#define INT_SIZE 4          // size of int data type in bytes
#define LOCKT_SIZE 4        // size of omp_lock_t data type in bytes

typedef int tVector[SIZE1];

typedef struct {
    int bin;
    int nelem;
    omp_lock_t lock;
} tBin;
typedef tBin tHisto[SIZE2];

void build_histogram (tVector v, tHisto h) {
    #pragma omp parallel
    #pragma omp single
    {
        for (int k=0; k<SIZE1; k++)
            #pragma omp task
            {
                int ind = v[k] % SIZE2;
                omp_set_lock(&h[ind].lock);
                if (h[ind].nelem < MAX) {
                    h[ind].bin += v[k];
                    h[ind].nelem++;
                }
                omp_unset_lock(&h[ind].lock);
            }
    }
}
```

Each entry of the histogram (`tHisto`) contains information about the values accumulated in each bin of the histogram (field `bin`) and the number of elements in each bin (field `nelem`). The program only accumulates the first `MAX` elements that go to each bin. The implementation makes use of locks (field `lock`) to protect the appropriate access to each element of type `tBin`; you can assume that locks are appropriately created and destroyed in the main program.

1. Assuming that the granularity for the tasks generated is not an issue, identify the two main issues that clearly degrade the performance during the execution of the parallel region in the histogram computation. Identify which parts of the code and/or data accesses are the responsible for that performance degradation.

2. Write a new version of the parallel code that follows **the same task decomposition** and solves the two previously identified performance issues.

3. The implementation of the `omp_set_lock` and `omp_unset_lock` function in our experimental system is as follows:

```
typedef int omp_lock_t;
// *lock is initialized to 0
void omp_set_lock (omp_lock_t *lock) {
    while (compare_and_swap(lock, 0, 1)==1);
}
void omp_unset_lock (omp_lock_t *lock) {
    *lock=0;
}
```

where the atomic function `compare_and_swap (omp_lock_t *p, int old, int new)` works as follows: it compares the memory location pointed by `p` with `old`, and if they are the same, the content of this memory location is set to `new`, otherwise is unchanged. The function returns the original value in the memory location pointed by `p`. **We ask you** to write an implementation for the atomic function `compare_and_swap` that makes use of the built-in *load-linked* and *store-conditional* operations. The *load-linked* operation is performed by a call to `load_linked(int *p)` and the *store-conditional* operation is performed by a call to `int store_cond(int *p, int value)`.

Problem 2 (4 points) The following code excerpt implements the multiplication of the lower triangular part of a dense matrix `T` times a vector `X`, accumulating the result in vector `Y`:

```
/* Y += T * X */
...
for (i=0; i<N; i++) {
    for (j=0; j<=i; j++) {
        Y[i] += T[i][j] * X[j];
    }
}
...
```

We want to parallelise the code using appropriate OpenMP pragmas and invocations to intrinsic functions. You can NOT make use of the `for` work-sharing construct in OpenMP to distribute work among threads. You can assume that the number of threads evenly divides `N`. **We ask you** to:

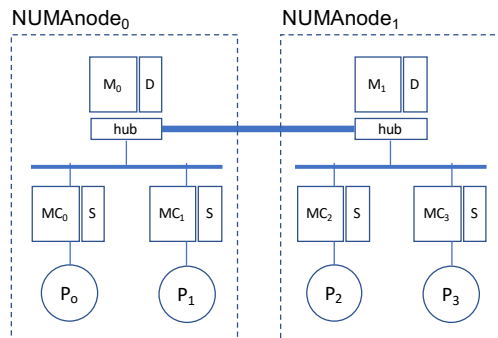
1. Implement a parallel version that obeys to an *input block geometric data decomposition* of the input vector `X`.

2. Implement a new parallel version that obeys to an *output block geometric data decomposition* of the output vector Y .

3. Improve the previous last parallel version by implementing an enhanced parallel version which improves load balancing. This new version should also exploit spatial data locality and avoid false sharing.

Note: If needed, you can assume that the number consecutive of elements of the data structures that fit into a cache line is `NUM_ELEMENTS_PER_CACHE_LINE`.

Problem 3 (3 points) Consider the following parallel architecture composed of two NUMA nodes, each with its own main memory, directory to keep coherence between NUMA nodes (write-invalidate MSU) and two processors with their own cache memory and snoopy to keep coherence within each node (write-invalidate MSI). For the purposes of this problem, you can assume: 1) infinite sizes for both main and cache memories; and 2) line size stores 16 consecutive integer values.



Legend:

NUMAnode_i: NUMA node i with two processors
 M_i : main memory for NUMAnode _{i}
 D : directory associated to M
hub: interconnect between NUMA nodes
 P_j : processor j inside NUMA node
 MC_j : cache memory local to processor P_j
 S : snoopy associated to MC

Coherence commands:

- Snoopy: BusRd _{j} , BusRdX _{j} and Flush _{j} , being j the snoopy/cache number doing the action
- Hub/directoty: RdReq _{$i \rightarrow j$} , WrReq _{$i \rightarrow j$} , Dreply _{$i \rightarrow j$} , Fetch _{$i \rightarrow j$} , Invalidate _{$i \rightarrow j$} and WriteBack _{$i \rightarrow j$} , from NUMAnode _{i} to NUMAnode _{j}


Also consider the following skeleton for a parallel program accessing to matrices a and b :

```
#define CACHE_LINE 16 // number of integers in a cache line
#define n 64
int a[n][n], b[n][n];
...
// initialization loop
for (int i=0; i<n; i++)
    for (int j=0; j<n; j++) {
        a[i][j] = ...;
        b[i][j] = ...;
    }

// compute loop 1: can only be parallelized with 2 processors, only dimension i
for (int i=0; i<n; i++)
    for (int j=0; j<n; j++)
        a[i][j] = foo(a[i][j], b[i][j]);


// compute loop 2: can be parallelized with 4 processors, both dimensions i and j
for (int i=0; i<n; i++)
    for (int j=0; j<n; j++)
        b[i][j] = goo(b[i][j], a[j][i]);
```

Functions foo and goo do not have any lateral effect on memory (i.e. don't modify any value in memory), variables i and j are stored in registers and elements of matrices are stored by rows. Assuming that after the initialization loop matrices a and b are mapped to NUMA nodes as follows with NO copies in cache (i.e. all of them in U state in the corresponding directory):




1. How many entries in the directory associated to each main memory M_0 and M_1 are occupied by matrices a and b ? How many bits are necessary in each directory entry to keep coherence information? Clearly indicate which information is stored in these bits.


2. Assume that only the i loop can be parallelized in *compute loop 1*, using processors P_0 and P_2 , as shown in the iteration space below:



During the execution of *compute loop 1*, which coherence commands are placed on the bus by the snoopies and which coherence commands are exchanged between NUMA nodes to maintain the coherence for the 4 memory lines of matrices a and b shown below? After the execution of *compute loop 1*, which is going to be the status for these 4 lines of matrices a and b both in the cache memories and in the directories? Fill in the table in the solutions page to answer this question.



3. Assume that both loops can be parallelized in *compute loop 2*, using all processors in the system as shown in the iteration space below:



During the execution of *compute loop 2*, which coherence commands are placed on the bus by the snoopies and which coherence commands are exchanged between NUMA nodes to maintain the coherence for the same 4 memory lines of matrices a and b? Please note the transposed access to matrix a in this loop. After the execution of *compute loop 2*, which is going to be the status for these 4 lines of matrices a and b both in the cache memories and in the directories? Fill in the table in the solutions page to answer this question.

Student name:

Table to be used to deliver your solution to **Problem 3**

		NUMA node 0					NUMA node 1				
		Bus command	MC0 status	MC1 status	NUMA commands	Directory entry	Bus command	MC2 status	MC3 status	NUMA commands	Directory entry
compute loop 1	line 0 for a										
	line 1 for a										
	line 2 for a										
	line 3 for a										
	line 0 for b										
	line 1 for b										
	line 2 for b										
	line 3 for b										
	compute loop 2	line 0 for a									
line 1 for a											
line 2 for a											
line 3 for a											
line 0 for b											
line 1 for b											
line 2 for b											
line 3 for b											